

## Front Matter

I want the Notebook to be as informative as possible, but model creating and training process follows some standard procedure that I do not want to repeat. Therefore, if you can, spend time reading the `PROLOGUE/Routine.ipynb` Notebook first.

## Paper Implementation - VGG16

Hello, this is my first milestone project - implementation of the VGG16 architecture from the paper "[Very Deep Convolutional Networks for Large-Scale Image Recognition](#)". The paper explored the effect of increasing layers on a model based on the filter size of  $3 * 3$  that we previously explored in the `TinyCNN` notebook (that is also the reason why that architecture is called TinyVGG). The architecture was the runner-up in the ImageNet 2014 Challenge for classification.

16 in VGG16 stands for 16 layers, where they are based on two basic units: convolution with filter size  $3 * 3$ , stride 1, padding 1 and max-pooling with window size  $2 * 2$ , stride 2. The table shown below, taken from the paper above, is the architecture for each of the VGG configuration. In this notebook, we will implement the VGG16-D one.

- 10.48550\_arxiv.1409.1556.pdf.png)

In this first notebook, we will focus on getting and transforming the data first.

## Downloading and extracting data

The task for our model will be classification, using a bigger dataset called [Food101](#). This is a built-in PyTorch dataset, so the processing can be fairly straightforward. However, it is not fun, so let's take the [Kaggle version](#) and process it to what we want.

First, downloading data from Kaggle. The easy way: you can download the zip file (~6 GB), upload it to Google Drive, and then mount Google Drive to Colab. . The slightly harder: you will need to sign up and obtain a Kaggle token, and then use the `kaggle` module to download the data. Let's do that.

In [ ]:

```
!pip install torchmetrics
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
```

```
Collecting torchmetrics
```

```
  Downloading torchmetrics-0.11.0-py3-none-any.whl (512 kB)
```

```
    |████████████████████████████████████████| 512 kB 4.3 MB/s
```

```
Requirement already satisfied: kaggle in /usr/local/lib/python3.8/dist-packages (1.5.12)
```

```
Requirement already satisfied: numpy>=1.17.2 in /usr/local/lib/python3.8/dist-packages (from torchmetrics) (1.21.6)
```

```
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.8/dist-packages (from torchmetrics) (4.4.0)
```

```
Requirement already satisfied: packaging in /usr/local/lib/python3.8/dist-packages
```

```
(from torchmetrics) (21.3)
Requirement already satisfied: torch>=1.8.1 in /usr/local/lib/python3.8/dist-packa
ges (from torchmetrics) (1.13.0+cu116)
Requirement already satisfied: certifi in /usr/local/lib/python3.8/dist-packages
(from kaggle) (2022.12.7)
Requirement already satisfied: tqdm in /usr/local/lib/python3.8/dist-packages (fro
m kaggle) (4.64.1)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.8/dist-pac
kages (from kaggle) (7.0.0)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.8/dist-packages
(from kaggle) (1.15.0)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.8/dist-packages
(from kaggle) (1.24.3)
Requirement already satisfied: requests in /usr/local/lib/python3.8/dist-packages
(from kaggle) (2.23.0)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.8/dist-pa
ckages (from kaggle) (2.8.2)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /usr/local/lib/python3.
8/dist-packages (from packaging->torchmetrics) (3.0.9)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.8/dis
t-packages (from python-slugify->kaggle) (1.3)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.8/dist-
packages (from requests->kaggle) (3.0.4)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.8/dist-packa
ges (from requests->kaggle) (2.10)
Installing collected packages: torchmetrics
```

In [ ]:

```
!pip install --upgrade mlxtend kaggle
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheel
s/public/simple/
Requirement already satisfied: mlxtend in /usr/local/lib/python3.8/dist-packages
(0.14.0)
Collecting mlxtend
  Downloading mlxtend-0.21.0-py2.py3-none-any.whl (1.3 MB)
    |████████████████████████████████████████| 1.3 MB 4.1 MB/s
Requirement already satisfied: kaggle in /usr/local/lib/python3.8/dist-packages
(1.5.12)
Requirement already satisfied: scipy>=1.2.1 in /usr/local/lib/python3.8/dist-packa
ges (from mlxtend) (1.7.3)
Requirement already satisfied: scikit-learn>=1.0.2 in /usr/local/lib/python3.8/dis
t-packages (from mlxtend) (1.0.2)
Requirement already satisfied: numpy>=1.16.2 in /usr/local/lib/python3.8/dist-pack
ages (from mlxtend) (1.21.6)
Requirement already satisfied: matplotlib>=3.0.0 in /usr/local/lib/python3.8/dist-
packages (from mlxtend) (3.2.2)
Requirement already satisfied: pandas>=0.24.2 in /usr/local/lib/python3.8/dist-pac
kages (from mlxtend) (1.3.5)
Requirement already satisfied: joblib>=0.13.2 in /usr/local/lib/python3.8/dist-pac
kages (from mlxtend) (1.2.0)
Requirement already satisfied: setuptools in /usr/local/lib/python3.8/dist-package
s (from mlxtend) (57.4.0)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.8/di
st-packages (from matplotlib>=3.0.0->mlxtend) (2.8.2)
Requirement already satisfied: pyparsing!=2.0.4,!2.1.2,!2.1.6,>=2.0.1 in /usr/lo
cal/lib/python3.8/dist-packages (from matplotlib>=3.0.0->mlxtend) (3.0.9)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.8/dist-
packages (from matplotlib>=3.0.0->mlxtend) (1.4.4)
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.8/dist-packa
ges (from matplotlib>=3.0.0->mlxtend) (0.11.0)
Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.8/dist-packa
ges (from pandas>=0.24.2->mlxtend) (2022.6)
```

```

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.8/dist-packages
(from python-dateutil>=2.1->matplotlib>=3.0.0->mlxtend) (1.15.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.8/dist-packages
(from scikit-learn>=1.0.2->mlxtend) (3.1.0)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.8/dist-packages
(from kaggle) (1.24.3)
Requirement already satisfied: requests in /usr/local/lib/python3.8/dist-packages
(from kaggle) (2.23.0)
Requirement already satisfied: certifi in /usr/local/lib/python3.8/dist-packages
(from kaggle) (2022.12.7)
Requirement already satisfied: tqdm in /usr/local/lib/python3.8/dist-packages (from
kaggle) (4.64.1)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.8/dist-packages
(from kaggle) (7.0.0)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.8/dist-packages
(from python-slugify->kaggle) (1.3)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.8/dist-packages
(from requests->kaggle) (2.10)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.8/dist-packages
(from requests->kaggle) (3.0.4)
Installing collected packages: mlxtend
  Attempting uninstall: mlxtend
    Found existing installation: mlxtend 0.14.0
    Uninstalling mlxtend-0.14.0:
      Successfully uninstalled mlxtend-0.14.0
Successfully installed mlxtend-0.21.0

```

First, we need a variable to keep track of the environment we are in as it is different to run the notebook right on Kaggle and run it anywhere else (from the teaching of a Kaggle Grandmaster)

```

In [ ]: # Import modules
import os
from pathlib import Path

# Keep an environment variable
iskaggle = os.environ.get('KAGGLE_KERNEL_RUN_TYPE', '')

```

Next, based on the [docs](#), we will need to create a `./kaggle/kaggle.json`. You can go to File Explorer and create a folder in your machine, or we can code that. I will code.

```

In [ ]: # Paste your API here (I have run and then deleted mine)
creds = ''

```

```

In [ ]: cred_path = Path('~/.kaggle/kaggle.json').expanduser()
if not cred_path.exists():
    cred_path.parent.mkdir(exist_ok=True)
    cred_path.write_text(creds)
    cred_path.chmod(0o600)

```

Next, let's use the method `dataset_download_cli` to download and unzip data files.

```

In [ ]: # Sanity check
!kaggle datasets list

```

```

ref                                     title
size  lastUpdated                    downloadCount  voteCount  usabilityRating
-----

```

meirnazri/covid19-dataset					COVID-19 Dataset
5MB	2022-11-13 15:47:17	11624	344	1.0	
michals22/coffee-dataset					Coffee dataset
24KB	2022-12-15 20:02:12	2316	63	1.0	
thedevastator/jobs-dataset-from-glassdoor					Salary Prediction
3MB	2022-11-16 13:52:31	7233	155	1.0	
thedevastator/unlock-profits-with-e-commerce-sales-data					E-Commerce Sales D
ataset	6MB 2022-12-03 09:27:17				1694 4
8 1.0					
ahmettalhabektas/argentina-car-prices					Argentina car pric
es	8KB 2022-12-05 09:05:21				745 3
4 1.0					
mvieira101/global-cost-of-living					Global Cost of Liv
ing	1MB 2022-12-03 16:37:53				3260 6
9 0.9705882					
thedevastator/uncovering-wage-disparities-in-pennsylvania-s-hi					Higher Education W
ages	223KB 2022-12-04 15:42:36				1172 3
6 1.0					
danela/fatal-alligator-attacks-us					Fatal Alligator At
tacks US	5KB 2022-12-15 16:37:57				350 2
6 1.0					
kabhishm/best-selling-music-artists-of-all-time					Best Selling Music
Artists of All Time	3KB 2022-12-09 07:04:29				920 39
0.9411765					
swaptr/fifa-world-cup-2022-statistics					FIFA World Cup 202
2 Team Data	15KB 2022-12-19 00:29:15				2478 5
8 0.9705882					
whenamancodes/predict-diabities					Predict Diabetes
9KB	2022-11-09 12:18:49	7549	119	1.0	
die9origephit/fifa-world-cup-2022-complete-dataset					Fifa World Cup 202
2: Complete Dataset	7KB 2022-12-18 22:51:11				1991 8
3 0.9411765					
mattp/alcohol-consumption-per-capita-2016					Alcohol Consumptio
n Per Capita 2016	4KB 2022-12-09 00:03:11				1186 4
3 1.0					
swaptr/fifa-world-cup-2022-match-data					FIFA World Cup 202
2 Match Data	7KB 2022-12-19 00:30:28				1201 3
3 1.0					
laibaanwer/superstore-sales-dataset					SuperStore Sales D
ataset	2MB 2022-12-07 08:53:32				1339 3
6 1.0					
thedevastator/discovering-hidden-trends-in-global-video-games					Discovering Hidden
Trends in Global Video Games	56KB 2022-12-03 11:21:47				727 39
1.0					
thedevastator/the-ultimate-netflix-tv-shows-and-movies-dataset					Netflix TV Shows a
nd Movies (2022 Updated)	2MB 2022-11-27 20:41:41				2272 3
9 1.0					
catherinerasgaitis/mxmh-survey-results					Music & Mental Hea
lth Survey Results	22KB 2022-11-21 10:03:12				2991 6
8 1.0					
kabhishm/imdb-100-movie-titles					IMDB 100 Movies
9KB	2022-12-07 11:36:06	655	32	0.9411765	
tirendazacademy/fifa-world-cup-2022-tweets					FIFA World Cup 202
2 Tweets	1MB 2022-12-08 19:43:37				958 3
1 1.0					

In [ ]:

```
path = Path('kmader/food41')
```

```
In [ ]: if not iskaggle and not path.exists():
        import kaggle
        kaggle.api.dataset_download_cli(str(path))
```

Downloading food41.zip to /content

100%|██████████| 5.30G/5.30G [02:58<00:00, 31.9MB/s]

We have the data in the zip file. Now all we need to do is to extract them out.

```
In [ ]: folder_path = Path('food41')
        os.mkdir(folder_path)
```

```
In [ ]: import zipfile
        zipfile.ZipFile(f'{folder_path}.zip').extractall(folder_path)
```

## Food-101

The dataset is introduced in this [paper](#), consisting of 101 classes, each with 750 training and 250 testing examples, totalling 1000 images each. The dataset comes with a metadata folder, giving information about which image should go into which subset, which is great! The data was already split into training and testing examples, but we also need a *validation set*. The testing examples were manually selected to contain noise and challenge the model, so we will not touch that, but we will split the training set further to create a validation set. Now, creating a good validation set is [an art](#), but here we will just use good ol' random splitting.

First, we will need to format the images folder into train and test folders. Next, we will load the data. The process is quite the same, what's new this time is we will random split the training data into training set and validation set, as well as applying more transformation.

```
In [ ]: # Generic torch process
        from torch import nn
        import torch
        from torch.utils.data import DataLoader

        # Specifically for computer vision
        import torchvision
        from torchvision import datasets, transforms

        # Other module(s)
        import matplotlib.pyplot as plt
        import gc
        import json
        import shutil
        import itertools
```

```
In [ ]: with open('/content/food41/meta/meta/train.json', 'r') as fp:
        train_dict = json.load(fp)
        with open('/content/food41/meta/meta/test.json', 'r') as fp:
            test_dict = json.load(fp)
        print(len(train_dict['apple_pie']), train_dict['apple_pie'][-10:])
        print(len(test_dict['apple_pie']), test_dict['apple_pie'][-10:])
```

```
750 ['apple_pie/960233', 'apple_pie/960669', 'apple_pie/962315', 'apple_pie/966595',
      'apple_pie/973088', 'apple_pie/973428', 'apple_pie/98352', 'apple_pie/98449', 'apple_pie/987860', 'apple_pie/997124']
250 ['apple_pie/885848', 'apple_pie/886793', 'apple_pie/904832', 'apple_pie/908367',
      'apple_pie/963140', 'apple_pie/981895', 'apple_pie/984571', 'apple_pie/986844', 'apple_pie/99556', 'apple_pie/997950']
```

The list value of a dictionary key contains the strings that are the file paths of the images without the extension. We will use this to copy the images to proper folders.

```
In [ ]: os.mkdir('data')
```

```
In [ ]: new_data_path = Path('data')
        original_data_path = Path('food41/images')
        new_folders = ['train', 'test']
        for folder in new_folders:
            if folder == 'train':
                for key, value in train_dict.items():
                    value_set = set(value)
                    if not os.path.exists(new_data_path/folder/key):
                        os.mkdir(new_data_path/folder/key)
                    for image in os.listdir(original_data_path/key):
                        image_path = key + '/' + image
                        image_path = image_path.split('.')[0]
                        if image_path in value_set:
                            shutil.copy(original_data_path/key/image, new_data_path/folder)
            else:
                for key, value in test_dict.items():
                    value_set = set(value)
                    if not os.path.exists(new_data_path/folder/key):
                        os.mkdir(new_data_path/folder/key)
                    for image in os.listdir(original_data_path/key):
                        image_path = key + '/' + image
                        image_path = image_path.split('.')[0]
                        if image_path in value_set:
                            shutil.copy(original_data_path/key/image, new_data_path/folder)
```

And we are done! Now we can load data as we like!

But first, let's write some transformations for the images to perform data augmentation.

## Data Augmentation

This is a technique to generate more training data by performing operations on the original data (such as flipping, shearing, rotating for images). The artificial data should generate the same output as the original one, but they are different, so hopefully the model is encouraged to learn the general pattern of the data instead of overfitting. Data augmentation is usually seen in training, but there is a technique called "test-time augmentation" that has been

passed down among the Kaggle Grandmaster and implemented in the library [fastai](#).

For the transformations, first we have the staples: `ToTensor()`, which turns images to `torch.tensor` objects. You may notice the `Normalize()` with some arbitrary parameters (the first is a list of means for each color channel and the second is a list of standard deviations for each color channel). These are parameters for the normalization of [ImageNet](#) dataset and are required by all PyTorch pre-trained models. This may not necessarily be true for our data, but it can be used. PyTorch also recommends having images of size  $224 * 224$  pixels, so we use `resize` to that. The other transformations do what it is called for, with parameters for angle, probability, etc. (Explore more transformations on PyTorch [docs](#).)

```
In [ ]: train_transforms = transforms.Compose([transforms.RandomResizedCrop(224),
                                             transforms.RandomRotation(35),
                                             transforms.RandomVerticalFlip(0.27),
                                             transforms.RandomHorizontalFlip(0.27),
                                             transforms.ToTensor(),
                                             transforms.Normalize([0.485, 0.456, 0.406],

valid_n_test_transforms = transforms.Compose([transforms.Resize(224),
                                             transforms.ToTensor(),
                                             transforms.Normalize([0.485, 0.456, 0.406],
```

```
In [ ]: data_dir = Path('data')
train_dir = data_dir/'train'
test_dir = data_dir/'test'
```

```
In [ ]: train_dataset = datasets.ImageFolder(train_dir, transform = train_transforms)
```